

Trioron 2.0: a multi-axis substrate for resource-bounded continual learning

Marcelinus R. Hatorangan*

May 22, 2026

Abstract

We describe *Trioron 2.0*, a substrate-level revision to the trioron architecture introduced in the companion paper, exposing six per-cell axes of structural plasticity: long-range input reach, plastic fanout sparsity, depth growth, slow axonal gain, dendritic compartmentalization, and field-conditional layer emergence. The six axes share one design principle — each is a separate write channel into the substrate, with its own provenance and its own plasticity gate, rather than a feature of the loss or the optimizer. A small spatial primitive (positions on a unit square, partitioned into pools) then enables a cell-granularity composition operator: two organisms trained independently on disjoint task slices can be absorbed into one substrate without retraining when they share the L0 perception layer’s deterministic factor (the R·S handshake). The cell-granularity composition is compatible with positional locality constraints (locally-connected masks) because absorption respects pools, not exact positions. We report the mechanism’s first verification on a chained-15 disjoint-class slice: task-aware accuracy is preserved on both donors within $\Delta \leq 0.026$, full-softmax accuracy lifts by +0.20 to +0.31 via manifold merge plus a 200-step head-only calibration pass, and the post-absorption substrate is *more* locality-clean than either donor alone. The 2.0 substrate is backwards-compatible with 1.0 donors: legacy checkpoints load unchanged and run under the extended forward path with the new axes at their default-off settings.

*Correspondence: marcelinusrocky@aol.com. Repository: <https://github.com/Marcrockhat/trioron-project>. Live demo: huggingface.co/spaces/Marcrockhat/trioron-demo. Guided tour: marcrockhat.github.io/trioron-project/tour.

1	Introduction
2	Related Work
3	Methods
3.1	The Trioron Node, recapitulated
3.2	Axis 1: Long-range input reach
3.3	Axis 2: Plastic fanout sparsity
3.4	Axis 3: Depth growth
3.5	Axis 4: Slow axonal gain
3.6	Axis 5: Dendritic compartmentalization
3.7	Axis 6: Field-conditional layer emergence
3.8	Spatial substrate: positions, pools, and locality
3.9	Pool-matched absorption
4	Results
4.1	Substrate self-tests: per-axis end-to-end
4.2	Pool-matched absorption: the LCN unlock
4.3	Backward compatibility
5	Discussion
6	Conclusion
A	Detailed Methods
A.1	Locally-Connected mask construction
A.2	Pool partition
A.3	Pool-matched absorption procedure
A.4	Manifold archive
B	Reproducibility